

# Cross-Language Component Testing: Performance and Interoperability Insights

Kai Erik Niermann

June 25, 2024

# Outline

- 1 Introduction
  - Motivation
- 2 Research Questions
- 3 Study design
  - System Design
  - Approach
- 4 Methodology
  - rewriting
  - choice of apis
  - testing setup
  - mapping
  - improved static mapping
- 5 Discussion
- 6 Results

# Introduction

- Establishing terminology
- **Rewrite:**
  - Writing new code
  - Same or new language
  - Implements existing features
- **Refactoring:**
  - Rearrange/minimize existing code
- Reasons for rewrites:
  - Second system effect
  - New developments in libraries/languages
  - Initial choices restrictive as application scales
- Benefits and drawbacks:
  - Overimplementation risk (Fred Brooks)
  - Better understanding and new technologies
- Rewrites in software development:
  - Key component
  - Thesis explores effective approaches

# Motivation

- **Definition:**

- Using a different language for rewriting existing code

- **Benefits:**

- Addresses fundamental limitations of the original language

- **Challenges:**

- Effectiveness of new language
- Large system size
- Unforeseen incompatibilities

- **Testing Approach:**

- Test a critical component first
- Minimal additional code

- **Strategies:**

- Wrap component in testing framework
- Use language interop APIs

- **Objective:**

- Determine if the component works better in another language

# Research Questions

## Main Question

How can we test the performance of a system component rewritten in another programming language?

## Role of Language APIs

How do language APIs differ in their approaches to mapping objects and types, are some more restrictive than others?

## Performance and Real world applications

How does the overhead and restrictions on optimization techniques effect the application of interop APIs in performance dependent production code?

# system design




Figure: Ray tracer overview

- Complex enough system to demonstrate the idea in an applied setting.
- Distinct components that can be isolated for testing.
- Requires high performance to function effectively, making it ideal for assessing performance needs.

# Testing Approach: Q & A Format

- How do we isolate the most performance impacting component?
  - Profile our application for runtime and memory.
  - Use the profiling data to make a judgement on the most performance impacting component.
- What languages do we choose for rewriting?
  - Choose languages with embedding APIs for Julia.
  - Opt for Python and C++ due to their differing fundamental properties and available interface libraries.
- How do we test the rewritten components?
  - Use the language specific benchmarking tools to test components in isolation.
  - Use Julia's benchmarking tools to test components & overhead

# rewriting

- Most performance impact: hit functions, crucial for ray tracing efficiency.
- Efficiency concern: If target languages perform poorly with hit functions, rewriting the entire ray tracer may not be feasible.
- Rewriting components: Trivial task due to similar syntax across languages.

## C++

```
bool hit(const aabb& box, const ray& r, const interval& ray_t);
```

## Python

```
def hit(bbox: aabb, r: ray, ray_t: interval[float]) -> bool:
```

## Julia

```
function hit!(bbox::aabb, r::ray, ray_t::interval)::Bool
```



# Choosing an Embedding API

Choosing the right embedding API to attach hit functions to the ray tracer implementation was crucial.

## Python

- Python/JuliaCall API:
  - Builds on previous implementations.
  - Supports extensive language conversions.

## C++

- Options:
  - CxxWrap: Feature-rich but less user-friendly.
  - **(choice)** Jluna: Newer, user-friendly, lacks some features

## Considerations

- 1 Choose an API that minimizes additional code.
- 2 Prefer an API that simplifies component attachment.

Languages supporting reflection reduce additional code, simplifying rewriting and testing.

# Testing setup

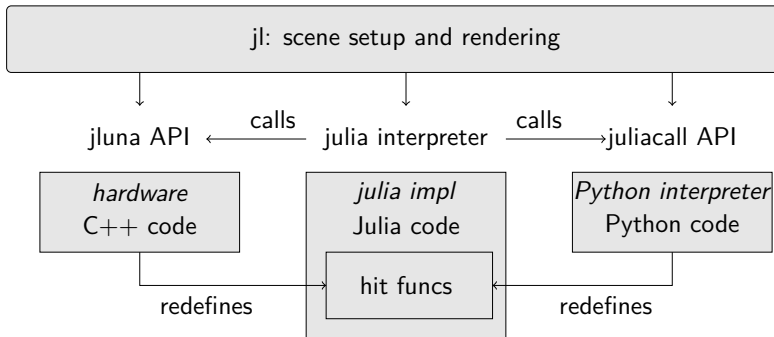


Figure: Testing setup for component isolation

# dynamic mapping

# Mapping Objects to Statically Typed Languages

## Static Mapping

- User describes interface for types (e.g., class/struct).
- API maps Julia type to corresponding C++ type.
- User-defined reflection system parses low-level types.

## Getters and Setters:

- Define functions  $\text{get}(x)$  and  $\text{put}(x)$ .
- Property  $x$  represented as a string.
- $\text{get} : S \rightarrow V = \text{object property } x$
- $\text{set} : S, V \rightarrow S = \text{modified property } x$

**Note:** Similar to functional programming lens.

# Challenges with Getter/Setter Pairs

## Limitation

Issues arise when source  $S$  has an abstract type.

## Notation

- $[S : T]$ ,  $[V : T]$  represent source and value types.
- $\text{set} : [S : T_B], [V : T_{bi}] \rightarrow S_m$

*Problem:* Julia uses strings for type representation:

- $\text{set} : [S : T_B], [V : \text{str}(T_{bi})] \rightarrow S_m$
- Need to map strings to C++ types.

*Solution:* Create a mapping mechanism

- string to type mapping  $st$  and type comparison  $f$ :

$$st : \text{str}(T_i) \rightarrow \text{typeid}(T_i) \quad f : st, \text{typeid}(T_p) \rightarrow \text{bool}$$

- Use compile-time code generation with variadic templates.

Achieve correspondence using fold expressions on variadic templates.

# Optimizing Setter Functions

An issue with the current approach:

- Requires providing all derived types to each getter/setter pair.
- Leads to significant overhead and potential errors.

Preferable solution:

- Define the set of derived types once.
- Allow all setters for attributes of a class to access this information.

# Improved Design with Varadic Templates

We can improve upon this design by utilizing varadic templates for types, getters, and setters to create a compile-time mapping:

```
def initialize_type[T, attr..., derived...]() -> void:
  fold => (attr,
    mapping.insert(attr::name,
      (
        // getter
        [instance: T]() { instance[T](attr::get(instance)) }
        // setter
        [instance: T, jl_val: jlval, jl_type_string: str] {
          fold(derived,
            if (f(st(jl_type_string), typeinfo(derived))):
              attr::set(instance, instance[derived](jl_val))
              matched = true
            )
          if (!matched):
            attr::set(instance, instance[T](jl_val))
        }
      )
    )
  )
```

# Template Metafunctions for Object Properties

Template metafunctions are used to define object properties:

```
template <typename Ot, typename Ft, const char* name>
struct Property {
    static constexpr const char* get_name() { return name; }
    static std::function<Ft(Ot&>> getter;
    static std::function<void(Ot&, Ft)> setter;
};
```

Example property declaration:

```
Usertype<bvh_node>::initialize_type(
    tl<
        Property<bvh_node, Hittable*, #left>,
        Property<bvh_node, Hittable*, #right>,
        Property<bvh_node, aabb, #bbox>
    >(),
    tl<Triangle, Sphere, bvh_node>());
```



# Limitations

## Considerations

- Users must specify class properties via lambdas.
  - Macros can simplify this process but its not ideal.
- All derived types must be passed to objects with base type attributes.
  - Tradeoff in library complexity and maintainability vs user defined types.
- Still requires additional user defined type specifications to allow for introspection

### key challenge

Find some means of having the base type interface discern the correct derived type to instantiate.

# Benefits

This system offers

- Simplified object deduction process
- Ability to map abstract objects
- Utilization of modern C++ principles for futureproofing
- Increased maintainability and extensibility
- Possibility to adapt this when Reflection TS becomes available in C++26

## Key Takeaway

It provides a more concise approach compared to traditional methods, enhancing ease of implementation and potential for future extensions.

# Insights - RQ2

## Observations

- Possible to implement a generic method for polymorphic object mapping.
- Offers JuliaCall API capabilities with C++ speed.
- Requires Julia data copying due to Julia C API limitations.

## Main takeaways

- Metaprogramming approaches effective for object reflection
- We can map to static languages with minimal additional boilerplate
- Mostly straightforward testing and integration across languages


# Baseline Performance analysis

## Timing method

- Call to `bvh_hit` from Julia
- Execution time in target languages (C++ and Python)

## Results

- Julia:
  - Many calls near 0ms (no hit)
  - Some calls around 0.0010ms (ray intersection and computation)
- Python and C++:
  - Similar call time distribution (measured from Julia)
  - Overhead will be clearer in overhead-exclusive graphs




# Isolated Performance - RQ3

## Results

- C++ outperforms Julia and Python in component runtime.
- Substantial overhead from Julia API (dynamic dispatch, type inference).
- Julia performs well with type stability, hindered by dynamic calls.

## Conclusion

- High-performance interop between different languages is challenging.
- Effective on a smaller scale, but less viable for very different languages.



# Results and Takeaways

- Different language APIs have varied integration approaches.
- Adapting APIs to specific use cases is generally manageable depending on the language
- Metaprogramming is effective for implementing generic, extensible interop libraries.
- Performance overhead mainly due to language incompatibilities.
- Mitigation possible with languages using similar execution methods.
- APIs work well for testing components.
- Careful consideration needed for using these APIs in performance-dependent production code.